

# Paramathics Interview Task II

Martin D. Pham

January 2021

## 1 Problem

Given a sparse lower triangular matrix  $L \in \mathbb{R}^{n \times n}$  and a vector  $b \in \mathbb{R}^n$ , we apply forward substitution to solve for the system:

$$Lx = b$$

Storing  $L$  in compressed column format allows a naive forward substitution implementation to skip iterating over zero-valued contributions given by the zero-valued entries of  $L$ .

We can further optimize the serial algorithm using the strategy shown in Sympiler. Steps are as follows: (1) construct adjacency graph  $DG_L = (V, E)$  of  $L$  representing dependencies between columns in triangular solve; (2) using the sparsity of  $b$ , i.e.  $\beta = \{i | b_i \neq 0\}$ , perform depth-first search on  $DG_L$  starting from  $\beta$  to determine  $Reach_L(\beta)$ ; (3) only the columns in  $Reach_L(\beta)$  contribute to the non-zero RHS entries and so all other columns, i.e.  $V \setminus Reach_L(\beta)$ , can be skipped during iteration.

No parallel optimizations were successful. Some naive attempts at using reduction clause with  $+$  operator to sum entries but updating for triangular solve became an issue regarding untangling loop dependency when working in CCS format. Seems like motivation for decoupled symbolic analysis?

## 2 Results

OS	Ubuntu 18.04.5 LTS
Memory	31.2GiB
Processor	Intel® Core™ i7-8550U CPU @ 1.80GHz × 8
Compiler	g++ (Ubuntu 7.5.0-3ubuntu1 18.04) 7.5.0

Table 1: Machine details (System 76 Galago Pro 2018 Model)

Linear system	Naive compressed column	Adjacency graph	Speedup
torso1	0.057025	0.034461	x1.6548
TSOPF	0.223503	0.223382	x1.0005

Table 2: Average wall time in seconds using *omp\_get\_wtime* over 10 runs between naive forward substitution compared to adjacency graph optimization.

Linear system	$ V $	$ Reach_L(\beta) $	$ Reach_L(\beta) / V $
torso1	116158	34314	0.30
TSOPF	35696	35414	0.99

Table 3: Total number of columns compared to number of contributing columns that need to be included in forward substitution. *torso1* skips 70% of columns while *TSOPF* only skips 1%.

The adjacency optimization improved performance for *torso1* but not *TSOPF*. Looking at the relative sizes of  $Reach_L(\beta)$  for both systems, we find that *torso1* is able to skip a significant number of columns while *TSOPF* barely benefits. The relative residuals in different norms also suggest that the matrix in *TSOPF* is ill-conditioned.

## References

Resources links attached to assignment.